

# Chapter 4: Deadlocks

Sukomal Pal, CSE, IIT(BHU)

# Chapter 4. Deadlocks

- *Deadlocks: Definition, Necessary and sufficient conditions for Deadlock, Deadlock Prevention*
- *Deadlock Avoidance: Banker's algorithm, Deadlock detection and Recovery.*

# INTRODUCTION

- *Deadlock* symbolizes a lock that is closed and whose key is, as if, lost.
- In OS it refers to a situation where a set of concurrent processes (or threads) perpetually block or starve for want of some resources held by some other processes (or threads) within the set.



**Fig 4.1:** Deadlock on road



**Fig 4.2:** Not a Deadlock

# INTRODUCTION

□ Deadlock is characterized by the following:

- it is caused for the want of *computing resources* (of any type).
- nature of *starvation* is *perpetual*.
- *starvation* occurs to *more than one* processes (or threads) simultaneously.
- the set of processes (or threads) have dependencies on each other in such a manner that they cannot come out of the perpetual stalemate on their own.

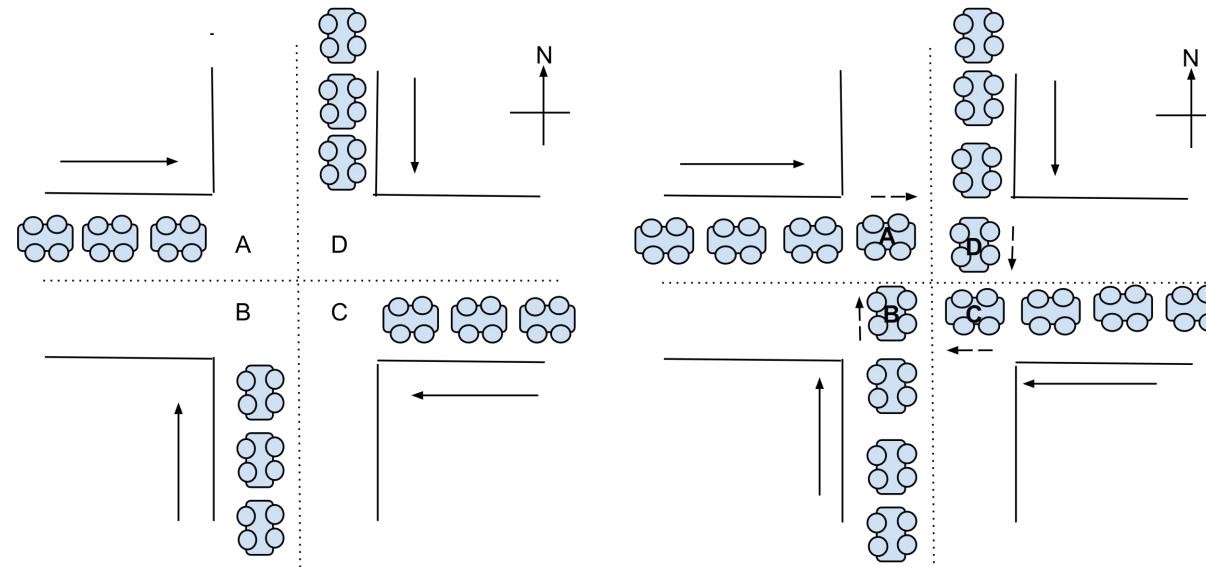


Fig 4.3a: Possibility of deadlock

Fig 4.3b: Deadlock

# Example of Deadlock

- **Livelock:** Continuous state of change without progress (e.g., two people trying to pass each other in a narrow hallway).
- Deadlock needs external intervention,livelock may resolve naturally.

## 1. Dining Philosophers:

- Each philosopher picks up their left fork first.
- All philosophers wait for the right fork indefinitely.
- Leads to **deadlock** as no philosopher can proceed.

## 2. Mutex Locks

- Two threads (Thread 1 & Thread 2) holding mutexes in different order.
- Each waits for the other's lock -> Deadlock occurs.

# Example of Deadlock

```
***** main process *****  
...  
pthread_mutex_t mutex1;  
pthread_mutex_t mutex2;  
...  
pthread_mutex_init(&mutex1, NULL);  
pthread_mutex_init(&mutex2, NULL);  
...  
  
***** thread1 *****  
void *function1(void *param)  
{  
    pthread_mutex_lock(&mutex1);  
    pthread_mutex_lock(&mutex2);  
    ...  
    /** * some work****/  
    ...  
    pthread_mutex_unlock(&mutex2);  
    pthread_mutex_unlock(&mutex1);  
    pthread_exit(0);  
}  
  
***** thread2 *****  
void *function2(void *param)  
{  
    pthread_mutex_lock(&mutex2);  
    pthread_mutex_lock(&mutex1);  
    ...  
    /** * some work****/  
    ...  
    pthread_mutex_unlock(&mutex1);  
    pthread_mutex_unlock(&mutex2);  
    pthread_exit(0);  
}
```

**Fig 4.4:** Possibility of a deadlock in a multi-threaded program

# Example of Deadlock

```
***** thread1 *****/
void *function1(void *param)
{
    while (1){
        pthread_mutex_lock(&mutex1);
        if (pthread_mutex_trylock(&mutex2)) {
            ...
            /* * some work*/
            ...
            pthread_mutex_unlock(&mutex2);
            pthread_mutex_unlock(&mutex1);
            break;
        }
        else
            pthread_mutex_unlock(&mutex1);
    } /* end of while */
    pthread_exit(0);
}
```

```
***** thread2 *****/
void *function2(void *param)
{
    while (1){
        pthread_mutex_lock(&mutex2);
        if (pthread_mutex_trylock(&mutex1)) {
            ...
            /* * some work*/
            ...
            pthread_mutex_unlock(&mutex1);
            pthread_mutex_unlock(&mutex2);
            break;
        }
        else
            pthread_mutex_unlock(&mutex2);
    } /* end of while */
    pthread_exit(0);
}
```

**Fig 4.5:** Possibility of a livelock in a multi-threaded program

# Resources

- A computing system can have one or more instances of each resource type, but only a finite number of instances.
- **Types of Resources:**
  - Hardware: Processors, memory, I/O devices.
  - Software: Files, semaphores, sockets.
- **Shareable vs. Non-Shareable Resources:**
  - Shareable: Read-only files (no deadlock risk).
  - Non-Shareable: Mutex locks, I/O devices (deadlock possible).
- **Reusable vs. Consumable Resources:**
  - Reusable: Can be used multiple times (e.g., CPU, memory, semaphores).
  - Consumable: Disappears after use (e.g., messages, signals).

# Resources

## Resource access

- During execution, a thread needs and uses several resources.
- The uses always obey the following sequence.
  - i. *Request*: A thread makes a request to the OS kernel for one or more instances of a resource. If an instance of the resource is not available, The thread waits (or blocks) till it acquires an instance of the resource.
  - ii. *Use*: Once acquired, it uses the instance of the resource non-shareably.
  - iii. *Release*: After the use, the thread returns the resource back to the kernel.

# Resource Allocation Graph

## □ Resource Allocation Process:

- Request: Thread requests a resource.
- Use: Thread holds and uses the resource.
- Release: Thread returns the resource after use.

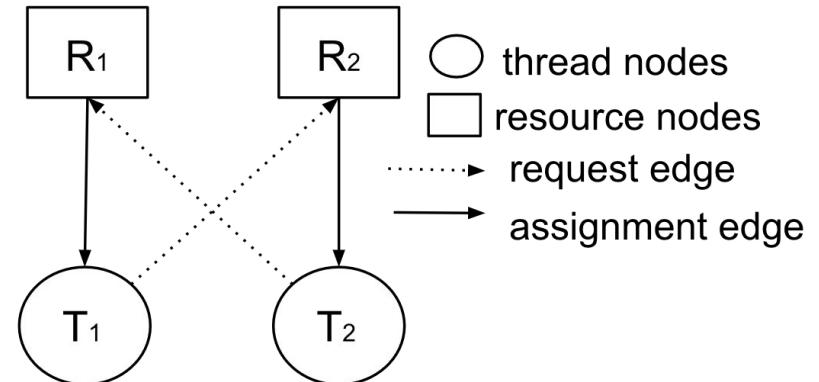
## □ Resource Allocation Graph (RAG):

- Nodes: Threads and Resources.
- Edges: Request edges ( $T \rightarrow R$ ) and Allocation edges ( $R \rightarrow T$ ).

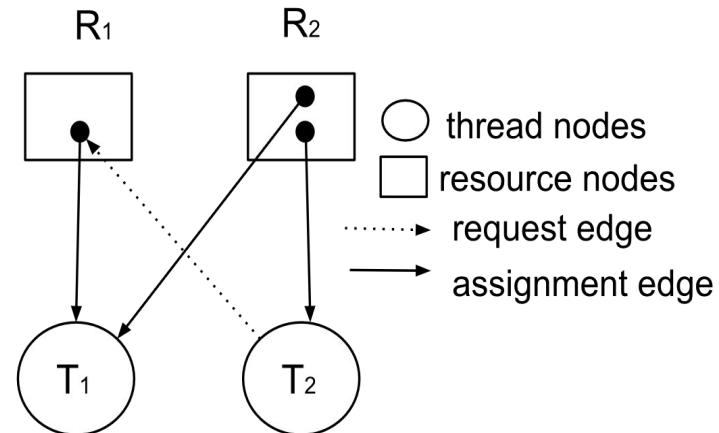
### □ Cycle in RAG = Deadlock.

## □ Multiple Instances of Resources:

- If multiple instances exist, a resource can be assigned to multiple threads.
- No cycle = No deadlock.



**Fig 4.4:** Resource Allocation Graph

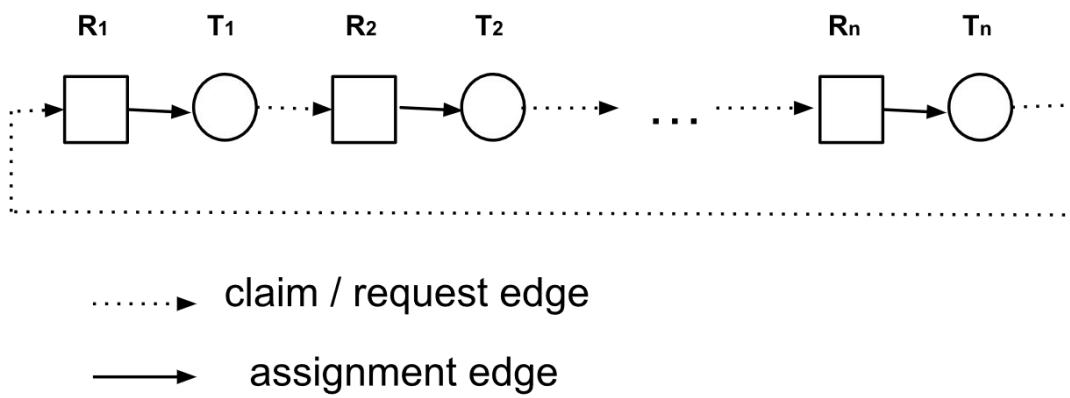


**Fig 4.5:** RAG With Multiple instances

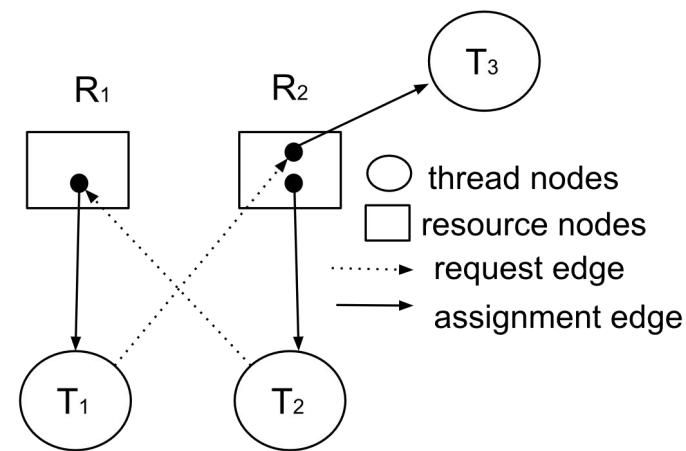
# CONDITIONS of A DEADLOCK

- i. **Mutual Exclusion:** When resources are used by threads non-sharably, then only there may emerge a possibility of deadlock.
  - There should be at least one resource that is used by threads in a mutually exclusive way.
- ii. **Hold and Wait:** During execution, threads are allowed to hold one or more resources and, at the same time, request to acquire a few more resources held by other thread(s).
- iii. **No Preemption:** None of the resources are preempted from the threads that hold them.
- iv. **Unresolvable Circular Wait:** A set of threads  $T = \{T_1, T_2, \dots, T_n\}$  hold and wait for resources from a set  $R = \{R_1, R_2, \dots, R_n\}$  in such a way that  $R_1 \rightarrow T_1, T_1 \rightarrow R_2, R_2 \rightarrow T_2, \dots, R_n \rightarrow T_n$  and  $T_n \rightarrow R_1$ .
  - Threads and resources make a cycle in the resource allocation graph with assignment and request edges.

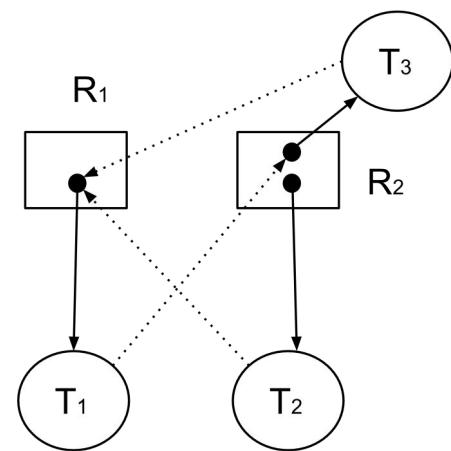
# Circular Wait



**Fig 4.6:** Circular Wait(RAG containing Cycle)



**Fig 4.7:** RAGs with multiple instances of resources



# HANDLING DEADLOCKS

- To stop occurrences of deadlocks, we must make sure that not all the four conditions are true at any point of time
- At least one of the four conditions must be negated

The strategies are clubbed into the following three categories.

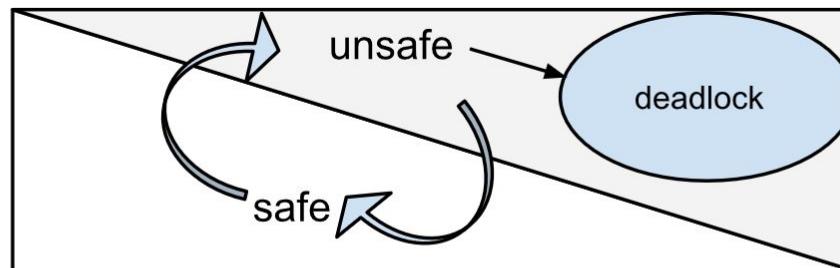
1. *Deadlock Prevention*: Requests to resources are monitored and allowed to be made only if all the four conditions are not satisfied simultaneously.
2. *Deadlock Avoidance*: The threads notify their overall need of resources in advance and the resources are allocated only if the allocation is *safe* (it does not lead to the possibility of a deadlock).
3. *Deadlock Detection & Recovery*: Deadlocks are allowed to happen. But they are detected, and appropriate recovery actions are taken.

# Deadlock Prevention

- 1. Preventing Mutual Exclusion**
- 2. Preventing Hold & Wait** : Ensure a process requests all resources at once.
- 3. Preventing No Preemption** : Allow resource reallocation from waiting processes.
- 4. Preventing Circular Wait** : Order resources numerically and request in ascending order.

# Deadlock Avoidance

- They allow the first three conditions (mutual exclusion, hold & wait, and no preemption) to continue.
- Safety of a system of threads is checked before any new allocation of resources.



**Fig 4.8:** Different states for a set of threads

# Banker's Algorithm

**Resources:** Total available resources are represented by a  $m$ -dimensional vector,

**Maximum resource needs:** Total requirement of all resource types by different threads is represented by a  $(n \times m)$  matrix

**Resource allocation:** Similarly, current allocation of resources at a given time, is also represented by another  $(n \times m)$  matrix,

**Available resources:** As resources are allocated to threads, free and available instances of resources reduce. The current number of available instances of resources is represented by an  $m$ -dimensional vector.

**Outstanding needs:** Once the threads are allocated resources, remaining resource needs of the threads are also represented by an  $n \times m$  matrix

**Resource requests:** Another matrix of  $(n \times m)$  dimension represents new (incremental) need of all the threads,

# Banker's Algorithm

The following relationships and constraints always hold true.

1.  $RES[j] \geq MAX[i][j]$  for all  $i, j$  (*maximum need of any thread for any resource-type cannot be more than the available number of instances in the system*)
2.  $RES[j] \geq \sum_i ALLOC[i][j]$  for all  $i, j$  (*sum of allocated instances of any resource-type cannot be more than total number of instances at any moment*)
3.  $AVAIL[j] = RES[j] - \sum_i ALLOC[i][j]$  for all  $i, j$  (*available number of resource-instances is whatever remains after allocation to all threads*)
4.  $NEED[i][j] = MAX[i][j] - ALLOC[i][j] \geq 0$  for all  $i, j$

# Banker's Algorithm

```
bool check_safety (AVAIL, ALLOC, NEED) {  
  
    0. INITIALISATION:  
        bool finish_possible[n] = {0, 0, ..., 0}; /*flag*/  
        bool safe = 1, unsafe = 0;  
        int WORK[m];  
        /* n = #threads, m=#types of resources */  
  
        WORK = AVAIL;  
  
    1. Find an index i such that  
        (NEED[i] < WORK)&&(finish_possible[i] == 0)  
        if not found, goto Step 3.  
  
    2. WORK = WORK + ALLOC[i];  
        finish_possible[i] = 1;  
        goto Step 1.  
  
    3. if (finish_possible[i] == 1 for all i) return (safe);  
        else return (unsafe);  
}
```

```
bool grant_request (AVAIL, ALLOC, REQ[i], MAX) {  
  
    0. bool grant_possible = 1, grant_not_possible = 0;  
  
    1. NEED[i] = MAX[i] - ALLOC[i];  
  
    2. if (REQ[i] > NEED[i])  
        return (error); /*request is more than max-need */  
  
    3. if (REQ[i] > AVAIL[i])  
        return (grant_not_possible); /* Ti must wait */  
  
    4. /* as if Ti is allocated the resources */  
        AVAIL = AVAIL - REQ[i];  
        ALLOC[i] = ALLOC[i] + REQ[i];  
        NEED[i] = NEED[i] - REQ[i];  
  
    5. if( check_safety(AVAIL, ALLOC, NEED) == safe))  
        return (grant_possible);  
        else return (grant_not_possible); /* Ti must wait */  
}
```

Fig 4.9: Banker's Algorithm

# Banker's Algorithm

□ Example: Consider the following snapshot of a system:

	<u>Allocation</u>				<u>Max</u>				<u>Available</u>			
	A	B	C	D	A	B	C	D	A	B	C	D
T0	0	0	1	2	0	0	1	2	1	5	2	0
T1	1	0	0	0	1	7	5	0				
T2	1	3	5	4	2	3	5	6				
T3	0	6	3	2	0	6	5	2				
T4	0	0	1	4	0	6	5	6				

Answer the following questions using the banker's algorithm:

- What is the content of the matrix **Need**?
- Is the system in a safe state?
- If a request from thread T1 arrives for (0,4,2,0), can the request be granted immediately?

# Deadlock Detection & Recovery

- OS detects deadlock and recovers from it, Detection based on **circular wait condition**

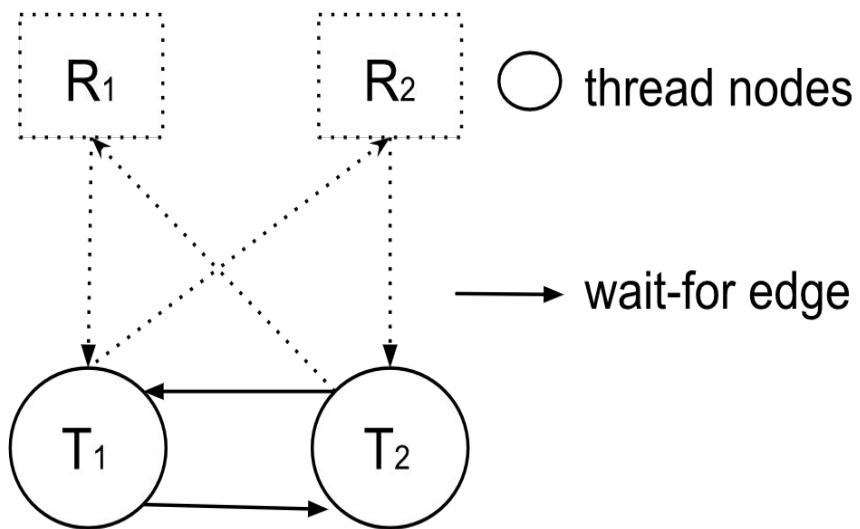
## **Case 1: Each resource has a single instance**

- Use a **Wait-For Graph (WFG)**
- A cycle in WFG **confirms deadlock**
- Cycle detection time complexity:  **$O(n^2)$**

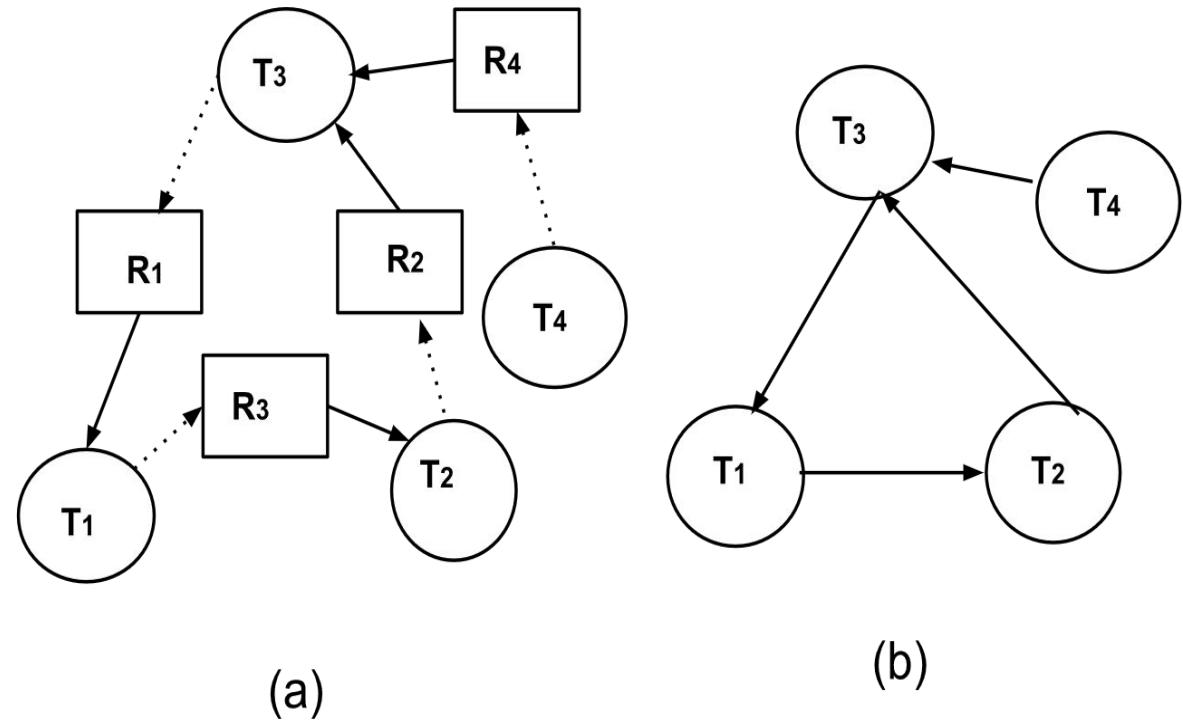
## **Case 2: Resources have multiple instances**

- Uses a **modified Banker's algorithm**
- Steps:
  - Threads with no allocated resources are ignored ( $O(n)$ )
  - Find a thread that can complete ( $O(mn)$ )
  - Simulate its completion and update availability ( $O(mn^2)$ )
  - If all threads finish, no deadlock; otherwise, deadlock exists
  - Time complexity:  **$O(mn^2)$**

# Deadlock Detection & Recovery



**Fig 4.10:** RAG(**Fig 4.4**) to Wait for Graph



**Fig 4.11:** (a) RAG to (b) wait for graph

# Deadlock Detection & Recovery

## **Resources with one or more instances:**

- When resources have multiple instances, mere presence of a cycle is not a confirmation of a deadlock.
- Use a deadlock detection algorithm

# Deadlock Detection & Recovery

```
bool detect_deadlock (AVAIL, ALLOC, REQ) {  
    0. INITIALISATION:  
        bool hold_res[n]; /*flags for threads*/  
        int WORK[m];  
            /* n = #threads, m=#types of resources */  
        WORK = AVAIL;  
  
    1. For all i, do  
        if (ALLOC[i] == [0, 0,...,0]) hold_res[i] = 0;  
        else hold_res[i] = 1;  
  
    2. Find an index i such that  
        (REQ[i] <= WORK) && (hold_res[i] == 1)  
        if not found, goto Step 4.  
  
    3. WORK = WORK + ALLOC[i];  
        ALLOC[i] = [0,0,...,0]; hold_res[i] = 0;  
        goto Step 2.  
  
    4. if (hold_res[i] == 1 for some i)  
        return (system in deadlock & Ti in deadlock);  
        else return (no_deadlock);  
}
```

**Fig 4.12:** Deadlock detection Algorithm

# Recovery from Deadlock

## Process Termination

- **Terminate all processes** in deadlock (brute-force, costly)
- **Terminate one process at a time** until deadlock is resolved
- Selection criteria:
  - Priority level (high, moderate, low)
  - Execution time (how long it has run)
  - Resource holding status (number of resources held)

# Recovery from Deadlock

## Resource Preemption

- **Forcibly reassign resources** from blocked threads to others
  - **Victim Selection:** which resources are to be chosen and from which threads?
  - **Rollback:** Restore the victim process to a safe state for future execution.
  - **Starvation Prevention:** Avoid repeated selection of the same victims.

# Reference

- [1] “OPERATING SYSTEMS”, Author: Dr. Sukomal Pal Associate Professor Department of Computer Science & Engineering IIT (BHU), Varanasi, UP